

Practical 4

Name: Meghna Dafade

PRN: 202501040044

4.1.1 Pandas – Series Creation and Manipulation

Problem Statement:

Write a Python program that takes a list of numbers from the user, creates a Pandas series from it, and then calculates the mean of even and odd numbers separately using the groupby and mean() operations.

Hint: You can group the Series based on whether each number is even or odd.

Input Format:

- The user should enter a list of numbers separated by space when prompted.

Output Format:

- The program should display the mean of even and odd numbers separately.
- Each mean value should be displayed with a label indicating whether it corresponds to even or odd numbers.

Code:

```
import pandas as pd
```

```
# Take inputs from the user to create a list of numbers
numbers = list(map(int, input().split()))
```

```
# Create a Pandas series from the list of numbers
series = pd.Series(numbers)
```

```
# Grouping by even and odd numbers and calculating the mean
grouped = series.groupby(series % 2 == 0).mean()
```

```
# Display the mean of even and odd numbers with labels
grouped.index = ['Even' if is_even else 'Odd' for is_even in grouped.index]
print("Mean of even and odd numbers:")
print(grouped)
```

Output:

seriesMa...

Submit

Average time
0.015 s
14.50 ms

Maximum time
0.034 s
34.00 ms

3 out of 3 shown test case(s) passed
3 out of 3 hidden test case(s) passed

Test case 1 34 ms Debug

Expected output
1 2 3 4 5 6 7 8 9 10
Mean of even and odd numbers:
Odd --- 5.0
Even --- 6.0
dtype: float64

Actual output
1 2 3 4 5 6 7 8 9 10
Mean of even and odd numbers:
Odd --- 5.0
Even --- 6.0
dtype: float64

Test case 2 15 ms Debug

Expected output
4 4 4 6 6 2 2 2 10 12 14 16
Mean of even and odd numbers:
Even --- 6.833333
dtype: float64

Actual output
4 4 4 6 6 2 2 2 10 12 14 16
Mean of even and odd numbers:
Even --- 6.833333
dtype: float64

Test case 3 10 ms Debug

Expected output
1 4 7 0 11 12 14 17 111 31 34
dtype: float64

Actual output
1 4 7 0 11 12 14 17 111 31 34
dtype: float64

Terminal

Test cases

4.1.2 Dictionary to Dataframe

Problem Statement:

A dictionary of lists has been provided to you in the editor. Create a DataFrame from the dictionary of lists and perform the listed operations, then display the DataFrame before and after each manipulation.

Create the DataFrame:

- Convert the dictionary to a Pandas DataFrame.

Add a new row:

- Take inputs from the user for the new row data (name, age).
- Add the new row to the DataFrame.
- Display the DataFrame after adding the new row

Modify a row:

- Modify a specific row by changing the age. Take the row index and new age value from the user.
- Display the DataFrame after modifying the row

Delete a row:

- Take the row index to be deleted from the user.
- Remove the specified row.
- Display the DataFrame after deleting the row

Add a new column:

- Add a column **Gender** with values taken from the user.
- Display the DataFrame after adding the new column

Modify a column:

- Convert names to uppercase.
- Display the DataFrame after modifying the column.

Delete a column:

- Remove the **Age** column.
- Display the DataFrame after deleting the column.

Code:

```
import pandas as pd
```

```
# Provided dictionary of lists
```

```
data = {  
    'Name': ['Alice', 'Bob', 'Charlie'],  
    'Age': [25, 30, 35],  
}
```

```
# Convert the dictionary to a DataFrame
```

```

df = pd.DataFrame(data)

# Display the original DataFrame
print("Original DataFrame:")
print(df)

# Adding a new
Name = input("New name: ")
Age = int(input("New age: "))
new = {'Name':Name, 'Age':Age}
df.loc[len(df)] = new
# Display the DataFrame after adding a new row
print("After adding a row:\n",df)

# Modifying a row
index = int(input("Index of row to modify: "))
age = int(input("New age: "))
df.at[index, 'Age']=age
# Display the DataFrame after modifying a row
print("After modifying a row:")
print(df)

# Deleting a row
delt = int(input("Index of row to delete: "))
df = df.drop(delt).reset_index(drop=True)
# Display the DataFrame after deleting a row
print("After deleting a row:")
print(df)

# Adding a new column
gen = input("Enter genders separated by space: ").split()
df['Gender']=gen
# Display the DataFrame after adding a new column
print("After adding a new column:")
print(df)

# Modifying a column
df['Name']=df['Name'].str.upper()
# Display the DataFrame after modifying a column
print("After modifying a column:")
print(df)

# Deleting a column

```

```
df = df.drop(columns = ['Age'])\
# Display the DataFrame after deleting a column
print("After deleting a column:")
print(df)
```

Output:

The screenshot displays a Jupyter Notebook interface with a 'Test cases' tab at the top. The status bar indicates '1 out of 1 shown test case(s) passed' and '1 out of 1 hidden test case(s) passed'. The notebook content shows a series of data manipulation steps on a DataFrame:

- Initial DataFrame:**

```

Name  Age
0  Alice  25
1    Bob  30
2  Charlie 35
```
- After adding a row:**

```

Name  Age
0  Alice  25
1    Bob  30
2  Charlie 35
3   Susan  28
```
- After modifying a row:**

```

Name  Age
0  Alice  27
1    Bob  30
2  Charlie 35
3   Susan  28
```
- After deleting a row:**

```

Name  Age
0  Alice  27
1    Bob  30
2  Charlie 35
```
- Enter genders separated by space:**

```

Name  Age  Gender
0  Alice  27      F
1    Bob  30      M
2  Charlie 35      M
```
- After adding a new column:**

```

Name  Age  Gender
0  Alice  27      F
1    Bob  30      M
2  Charlie 35      M
```
- After modifying a column:**

```

Name  Age  Gender
0  ALICE  27      F
1    BOB  30      M
2  CHARLIE 35      M
```
- After deleting a column:**

```

Name  Gender
0  ALICE      F
1    BOB      M
2  CHARLIE    M
```

4.1.3 Student Information

Problem Statement:

Write a program to read a text file containing student information (name, age, and grade) using Pandas. Perform the following tasks:

- Display the first five rows of the data frame.
- Calculate the average age of the students (limit the average age up to 2 decimal places).
- Filter out the students who have a grade above a certain threshold (consider the threshold grade is 'B').

Code:

```
import pandas as pd
```

```
# Read the text file into a DataFrame
```

```
file = input()
```

```
data = pd.read_csv(file, sep="\s+", header=None, names=["Name", "Age", "Grade"])
```

```
# write your code here..
```

```
print("First five rows:")
```

```
print(data.head())
```

```
average_age = round(data["Age"].mean(), 2)
```

```
print(f"Average age: {average_age}")
```

```
filtered_students = data[data["Grade"] <= "B"]
```

```
print("Students with a grade up to B")
```

```
print(filtered_students)
```

Output:

The screenshot shows a code execution environment with the following details:

- Performance Metrics:**
 - Average time: 0.080 s (79.50 ms)
 - Maximum time: 0.086 s (86.00 ms)
- Test Results:**
 - 1 out of 1 shown test case(s) passed
 - 1 out of 1 hidden test case(s) passed
- Test Case 1 (88 ms):**
 - Expected output:**

```
studentdata.txt
First five rows:
-- Name-- Age-- Grade
0-- John-- 25-- A
1-- Alin-- 22-- B
2-- Emma-- 24-- A
3-- Mike-- 23-- C
4-- Sony-- 21-- B
Average age: 23.29
Students with a grade up to B:
-- Name-- Age-- Grade
0-- John-- 25-- A
1-- Alin-- 22-- B
2-- Emma-- 24-- A
4-- Sony-- 21-- B
5-- Amia-- 26-- A
```
 - Actual output:**

```
studentdata.txt
First five rows:
-- Name-- Age-- Grade
0-- John-- 25-- A
1-- Alin-- 22-- B
2-- Emma-- 24-- A
3-- Mike-- 23-- C
4-- Sony-- 21-- B
Average age: 23.29
Students with a grade up to B:
-- Name-- Age-- Grade
0-- John-- 25-- A
1-- Alin-- 22-- B
2-- Emma-- 24-- A
4-- Sony-- 21-- B
5-- Amia-- 26-- A
```

4.2.1 Month with the Highest Total Sales

Problem Statement:

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the columns: Date, Product, Quantity, Price, and City.
- Group the data by Month and calculate the total sales for each month.
- Find the month with the highest total sales and display it.
- Also, display the total sales for the best month.

Code:

```
import pandas as pd
# Prompt the user for the file name
file_name = input()
# Load the data
df = pd.read_csv(file_name)
df['Date'] = pd.to_datetime(df['Date'])
df['Total_Sales_Amount'] = df['Quantity'] * df['Price']
Monthly_Sales = df.groupby(df['Date'].dt.strftime('%Y-%m'))['Total_Sales_Amount'].sum()

# Find the month with the highest total sales
best_month = Monthly_Sales.idxmax()
highest_sales = Monthly_Sales.max()
print(f"Best month: {best_month}")
print(f"Total sales: ${highest_sales:.2f}")
```

Output

The screenshot displays a code execution interface. At the top, performance metrics are shown: Average time 0.035 s (35.00 ms) and Maximum time 0.065 s (65.00 ms). Test results indicate 1 out of 1 shown test case(s) passed and 2 out of 2 hidden test case(s) passed. Below this, 'Test case 1' is detailed with a 65 ms execution time and a 'Debug' button. It compares 'Expected output' and 'Actual output' for three items: 'sales_data.csv', 'Best month: 2025-01', and 'Total sales: \$1210.00'. Both expected and actual outputs match for all items. At the bottom, there are tabs for 'Terminal' and 'Test cases'.

Metric	Value
Average time	0.035 s (35.00 ms)
Maximum time	0.065 s (65.00 ms)

Test Results:

- 1 out of 1 shown test case(s) passed
- 2 out of 2 hidden test case(s) passed

Test case 1 (65 ms) [Debug]

Expected output	Actual output
sales_data.csv	sales_data.csv
Best month: 2025-01	Best month: 2025-01
Total sales: \$1210.00	Total sales: \$1210.00

Terminal | Test cases

4.2.2 Best Selling Product

Problem Statement:

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:


- The CSV file contains the columns: Date, Product, Quantity, Price, and City.
- Find the product that sold the most in terms of quantity sold.
- Display the product that sold the most and the total quantity sold for that product.

Code:

```
import pandas as pd

# Prompt the user for the file name
file_name = input()
# Load the data
df = pd.read_csv(file_name)
product_sales = df.groupby("Product")["Quantity"].sum()
# Find the product with the highest total quantity sold
best_product = product_sales.idxmax()
highest_quantity = product_sales.max()
# Display the result
print(f"Best selling product: {best_product}")
print(f"Total quantity sold: {highest_quantity}")
```

Output:



The screenshot displays a code execution interface with the following components:

- Performance Metrics:**
 - Average time: 0.051 s (51.33 ms)
 - Maximum time: 0.094 s (94.00 ms)
- Test Results:**
 - 1 out of 1 shown test case(s) passed
 - 2 out of 2 hidden test case(s) passed
- Test Case Details (Test case 1, 94 ms):**
 - Expected output:**
 - sales_data.csv
 - Best selling product: Product A
 - Total quantity sold: 23
 - Actual output:**
 - sales_data.csv
 - Best selling product: Product A
 - Total quantity sold: 23
- Navigation:** Buttons for Terminal and Test cases are visible at the bottom.

4.2.3 City that Sold the Most Products

Problem Statement:

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the columns: Date, Product, Quantity, Price, and City.
- Group the data by City and calculate the total quantity of products sold for each city.
- Find the city that sold the most products (based on the total quantity sold).

Code:

```
import pandas as pd

# Prompt the user for the file name
file_name = input()
# Load the data
df = pd.read_csv(file_name)
# write the code..
city_sales = df.groupby("City")["Quantity"].sum()
best_city = city_sales.idxmax()

# Display the result
print(f"City sold the most products: {best_city}")
```

Output:

The screenshot displays a code execution interface. At the top, performance metrics are shown: Average time 0.025 s (25.00 ms) and Maximum time 0.045 s (45.00 ms). Test results indicate that 1 out of 1 shown test case(s) passed and 2 out of 2 hidden test case(s) passed. Below this, 'Test case 1' is detailed with a 45 ms execution time and a 'Debug' button. It compares 'Expected output' and 'Actual output', both showing 'sales_data.csv' and 'City sold the most products: Los Angeles'. At the bottom, there are tabs for 'Terminal' and 'Test cases'.

Metric	Value
Average time	0.025 s (25.00 ms)
Maximum time	0.045 s (45.00 ms)

Test Results:

- 1 out of 1 shown test case(s) passed
- 2 out of 2 hidden test case(s) passed

Test case 1 (45 ms) [Debug]

Expected output	Actual output
sales_data.csv	sales_data.csv
City sold the most products: Los Angeles	City sold the most products: Los Angeles

Terminal | Test cases

4.2.4 Most Frequently Sold Product Pairs

Problem Statement:

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the following columns: Date, Product, Quantity, Price, and City.
- For each date, find all pairs of products that were sold together (i.e., two products sold on the same date).
- Output the product pair/s that was sold most frequently.

Code:

```
import pandas as pd

from itertools import combinations
from collections import Counter

# Prompt user to input the file name
file_name = input()

# Read data from the specified CSV file
df = pd.read_csv(file_name)

# write the code
grouped = df.groupby("Date")["Product"].apply(list)
```

```

pairs_counter = Counter()
for products in grouped:
    pairs = combinations(sorted(products), 2)
    pairs_counter.update(pairs)

max_frequency = max(pairs_counter.values())
most_common_pairs = [(pair, freq) for pair, freq in pairs_counter.items() if freq ==
max_frequency]

# Output the most frequent product pairs
for(product1, product2), frequency in most_common_pairs:
    print(f"{product1} and {product2}: {frequency} times")

```

Output:

Average time
0.038 s
38.00 ms

Maximum time
0.076 s
76.00 ms

✓ 1 out of 1 shown test case(s) passed
✓ 2 out of 2 hidden test case(s) passed

✓ Test case 1 76 ms Debug

Expected output

Actual output

sales_data.csv	sales_data.csv
Product A and Product B: 2 times	Product A and Product B: 2 times
Product A and Product C: 2 times	Product A and Product C: 2 times

Terminal Test cases

4.2.5 Titanic Dataset Analysis and Data Cleaning

Problem Statement:

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset. For each question, perform necessary data cleaning, transformations, and calculations as required.

1. Display the first 5 rows of the dataset.
2. Display the last 5 rows of the dataset.
3. Get the shape of the dataset (number of rows and columns).
4. Get a summary of the dataset (using `.info()`).
5. Get basic statistics (mean, standard deviation, etc.) of the dataset using `.describe()`.
6. Check for missing values and display the count of missing values for each column.
7. Fill missing values in the 'Age' column with the median age.
8. Fill missing values in the 'Embarked' column with the most frequent value (`mode()`).
9. Drop the 'Cabin' column due to many missing values.
10. Create a new column, 'FamilySize' by adding the 'SibSp' and 'Parch' columns.

Code:

```
import pandas as pd
import numpy as np

# Load the Titanic dataset
data = pd.read_csv('Titanic-Dataset.csv')
print(data.head())

# 1. Display the first 5 rows of the dataset
print(data.tail())

# 2. Display the last 5 rows of the dataset
print(data.shape)

# 3. Get the shape of the dataset
print(data.info())

# 4. Get a summary of the dataset (info)
print(data.describe())

# 5. Get basic statistics of the dataset
print(data.isnull().sum())
```

6. Check for missing value

```
median_age = data['Age'].median()
```

```
data['Age'].fillna(median_age, inplace = True)
```

7. Fill missing values in the 'Age' column with the median age

```
mode_embarked = data['Embarked'].mode() [0]
```

```
data['Embarked'].fillna(mode_embarked, inplace = True)
```

8. Fill missing values in the 'Embarked' column with the mode

```
data.drop('Cabin', axis=1, inplace = True)
```

9. Drop the 'Cabin' column due to many missing values

```
data['FamilySize'] = data['SibSp'] + data['Parch']
```

10. Create a new column 'FamilySize' by adding 'SibSp' and 'Parch'

Output:

Average time: 0.223 s
Maximum time: 0.223 s
1 out of 1 shown test case(s) passed

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	FamilySize
889	1	1	30.00	889	890	1	1	30.00	889	890	1
C148	C			C148	C						
890	0	3	7.75	890	891	0	3	7.75	890	891	0
NaN	Q			NaN	Q						

[5 rows x 12 columns]
(891, 12)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	PassengerId	891 non-null	int64
1	Survived	891 non-null	int64
2	Pclass	891 non-null	int64
3	Name	891 non-null	object
4	Sex	891 non-null	object
5	Age	714 non-null	float64
6	SibSp	891 non-null	int64
7	Parch	891 non-null	int64
8	Ticket	891 non-null	object
9	Fare	891 non-null	float64
10	Cabin	204 non-null	object

< Prev Reset Submit Ne

Explorer

Average time
0.223 s
223.00 ms

Maximum time
0.223 s
223.00 ms

1 out of 1 shown test case(s) passed

50%
--- 0.000000 --- 0.000000 --- 3.000000 ---
--- 0.000000 --- 0.000000 --- 14.454200

75%
--- 1.000000 --- 0.000000 --- 31.000000
--- 1.000000 --- 0.000000 --- 31.000000
max
--- 8.000000 --- 6.000000 --- 512.329200

[8 rows x 7 columns]

PassengerId

Survived

Pclass

Name

Sex

Age

SibSp

Parch

Ticket

Fare

Cabin

Embarked

dtype: int64

50%
--- 0.000000 --- 0.000000 --- 3.000000 ---
--- 0.000000 --- 0.000000 --- 14.454200

75%
--- 1.000000 --- 0.000000 --- 31.000000
--- 1.000000 --- 0.000000 --- 31.000000
max
--- 8.000000 --- 6.000000 --- 512.329200

[8 rows x 7 columns]

PassengerId

Survived

Pclass

Name

Sex

Age

SibSp

Parch

Ticket

Fare

Cabin

Embarked

dtype: int64

Average time
0.223 s
223.00 ms

Maximum time
0.223 s
223.00 ms

1 out of 1 shown test case(s) passed

Test case 1 223 ms Debug

Expected output

Actual output

--- PassengerId --- Survived --- Pclass --- Fare --- Cabin --- Embarked

0 --- 1 --- 0 --- 3 --- 7.2500 --- NaN --- S

1 --- 2 --- 1 --- 1 --- 71.2833 --- C85 --- C

2 --- 3 --- 1 --- 3 --- 7.9250 --- NaN --- S

3 --- 4 --- 1 --- 1 --- 53.1000 --- C123 --- S

4 --- 5 --- 0 --- 3 --- 8.0500 --- NaN --- S

[5 rows x 12 columns]

--- PassengerId --- Survived --- Pclass --- Fare --- Cabin --- Embarked

886 --- 887 --- 0 --- 2 --- 13.00 --- NaN --- S

887 --- 888 --- 1 --- 1 --- 30.00 --- 842 --- S

888 --- 889 --- 0 --- 3 --- 23.45 --- NaN --- S

--- PassengerId --- Survived --- Pclass --- Fare --- Cabin --- Embarked

886 --- 887 --- 0 --- 2 --- 13.00 --- NaN --- S

887 --- 888 --- 1 --- 1 --- 30.00 --- 842 --- S

888 --- 889 --- 0 --- 3 --- 23.45 --- NaN --- S

4.2.6 Titanic Dataset Analysis and Data Cleaning-2

Problem Statement:

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset.

1. Create a new column 'IsAlone' which is 1 if the passenger is alone (FamilySize = 0), otherwise 0.
2. Convert the 'Sex' column to numeric values (male: 0, female: 1).
3. One-hot encode the 'Embarked' column, dropping the first category.
4. Get the mean age of passengers.
5. Get the median fare of passengers.
6. Get the number of passengers by class.
7. Get the number of passengers by gender.
8. Get the number of passengers by survival status.
9. Calculate the survival rate of passengers.
10. Calculate the survival rate by gender.

Code:

```
import pandas as pd
import numpy as np

# Load the Titanic dataset
data = pd.read_csv('Titanic-Dataset.csv')
data['FamilySize'] = data['SibSp'] + data['Parch']

# 1. Create a new column 'IsAlone' (1 if alone, 0 otherwise)
data['Alone'] = np.where(data['FamilySize'] == 0, 1, 0)

# 2. Convert 'Sex' to numeric (male: 0, female: 1)
data['Sex'] = data['Sex'].map({'male': 0, 'female': 1})

# 3. One-hot encode the 'Embarked' column
data = pd.get_dummies(data, columns = ['Embarked'], drop_first = True)

# 4. Get the mean age of passengers
mean_age = data['Age'].mean()
print(mean_age)

# 5. Get the median fare of passengers
median_fare = data['Fare'].median()
```

```
print(median_fare)
```

```
# 6. Get the number of passengers by class
```

```
passengers_by_class = data['Pclass'].value_counts()
```

```
print(passengers_by_class)
```

```
# 7. Get the number of passengers by gender
```

```
passengers_by_gender = data['Sex'].value_counts().sort_index()
```

```
print(passengers_by_gender)
```

```
# 8. Get the number of passengers by survival status
```

```
passengers_by_survival = data['Survived'].value_counts().sort_index()
```

```
print(passengers_by_survival)
```

```
# 9. Calculate the survival rate
```

```
survival_rate = data['Survived'].mean()
```

```
print(survival_rate)
```

```
# 10. Calculate the survival rate by gender
```

```
survival_rate_by_gender = data.groupby('Sex')['Survived'].mean()
```

```
print(survival_rate_by_gender)
```

Output:

Average time
0.108 s
108.00 ms

Maximum time
0.108 s
108.00 ms

1 out of 1 shown test case(s) passed

Test case 1 108 ms Debug

Expected output

Actual output

29.69911764705882

29.69911764705882

14.4542

14.4542

3 --- 491

3 --- 491

1 --- 216

1 --- 216

2 --- 184

2 --- 184

Name: Pclass, dtype: int64

Name: Pclass, dtype: int64

0 --- 577

0 --- 577

1 --- 314

1 --- 314

Name: Sex, dtype: int64

Name: Sex, dtype: int64

0 --- 549

0 --- 549

1 --- 342

1 --- 342

Name: Survived, dtype: int64

Name: Survived, dtype: int64

0.3838383838383838

0.3838383838383838

Sex

Sex

0 --- 0.188908

0 --- 0.188908

1 --- 0.742038

1 --- 0.742038

Name: Survived, dtype: float64

Name: Survived, dtype: float64

4.2.7 Titanic Dataset Analysis and Data Cleaning-3

Problem Statement:

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset.

1. Calculate the survival rate by class.
2. Calculate the survival rate by embarkation location (Embarked_S).
3. Calculate the survival rate by family size (FamilySize).
4. Calculate the survival rate by being alone (IsAlone).
5. Get the average fare by passenger class (Pclass).
6. Get the average age by passenger class (Pclass).
7. Get the average age by survival status (Survived).
8. Get the average fare by survival status (Survived).
9. Get the number of survivors by class (Pclass).
10. Get the number of non-survivors by class (Pclass).

Code:

```
import pandas as pd
import numpy as np

# Load the Titanic dataset
data = pd.read_csv('Titanic-Dataset.csv')
data['FamilySize'] = data['SibSp'] + data['Parch']
data['IsAlone'] = np.where(data['FamilySize'] > 0, 0, 1)
data = pd.get_dummies(data, columns=['Embarked'], drop_first=True)

# 1. Calculate the survival rate by class
print(data.groupby('Pclass')['Survived'].mean())

# 2. Calculate the survival rate by embarked location
print(data.groupby('Embarked_S')['Survived'].mean())

# 3. Calculate the survival rate by family size
print(data.groupby('FamilySize')['Survived'].mean())

# 4. Calculate the survival rate by being alone
print(data.groupby('IsAlone')['Survived'].mean())

# 5. Get the average fare by class
print(data.groupby('Pclass')['Fare'].mean())

# 6. Get the average age by class
```

```
print(data.groupby('Pclass')['Age'].mean())
```

7. Get the average age by survival status

```
print(data.groupby('Survived')['Age'].mean())
```

8. Get the average fare by survival status

```
print(data.groupby('Survived')['Fare'].mean())
```

9. Get the number of survivors by class

```
print(data[data['Survived'] == 1]['Pclass'].value_counts())
```

10. Get the number of non-survivors by class

```
print(data[data['Survived'] == 0]['Pclass'].value_counts())
```

Output:

The image displays two screenshots of a test runner interface, likely from a platform like LeetCode or a similar coding environment. Both screenshots show a test case passing with an average time of 0.108 s and a maximum time of 108.00 ms. The test case is labeled 'Test case 1' and is marked as '1 out of 1 shown test case(s) passed'.

The first screenshot shows the test results for a test case where the expected output and actual output are compared. The expected output is a dictionary with keys 'Pclass', 'Survived', 'Embarked', 'FamilySize', and 'Age'. The actual output is a dictionary with the same keys and values. The test case is marked as '1 out of 1 shown test case(s) passed'.

The second screenshot shows the test results for a test case where the expected output and actual output are compared. The expected output is a dictionary with keys 'Fare', 'Pclass', 'Age', 'Survived', and 'Fare'. The actual output is a dictionary with the same keys and values. The test case is marked as '1 out of 1 shown test case(s) passed'.

4.2.8 Titanic Dataset Analysis and Data Cleaning-4

Problem Statement:

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset.

1. Get the number of survivors by gender (Sex).
2. Get the number of non-survivors by gender (Sex).
3. Get the number of survivors by embarkation location (Embarked_S).
4. Get the number of non-survivors by embarkation location (Embarked_S).
5. Calculate the percentage of children (Age < 18) who survived.
6. Calculate the percentage of adults (Age >= 18) who survived.
7. Get the median age of survivors.
8. Get the median age of non-survivors.
9. Get the median fare of survivors.
10. Get the median fare of non-survivors.

Code:

```
import pandas as pd
import numpy as np
```

```
# Load the Titanic dataset
```

```
data = pd.read_csv('Titanic-Dataset.csv')
```

```
data = pd.get_dummies(data, columns=['Embarked'], drop_first=True)
```

```
# 1. Get the number of survivors by gender
```

```
survivors_by_gender = data[data['Survived'] == 1]['Sex'].value_counts()
print(survivors_by_gender)
```

```
# 2. Get the number of non-survivors by gender
```

```
non_survivors_by_gender = data[data['Survived'] == 0]['Sex'].value_counts()
print(non_survivors_by_gender)
```

```
# 3. Get the number of survivors by embarked location
```

```
survivors_by_embarked_s = data[data['Survived'] == 1]['Embarked_S'].value_counts()
print(survivors_by_embarked_s)
```

```
# 4. Get the number of non-survivors by embarked location
```

```
non_survivors_by_embarked_s = data[data['Survived'] == 0]['Embarked_S'].value_counts()
print(non_survivors_by_embarked_s)
```

5. Calculate the percentage of children (Age < 18) who survived

```
children = data[data['Age'] < 18]
children_survival_rate = children['Survived'].mean()
print(children_survival_rate)
```

6. Calculate the percentage of adults (Age >= 18) who survived

```
adults = data[data['Age'] >= 18]
adults_survival_rate = adults['Survived'].mean()
print(adults_survival_rate)
```

7. Get the median age of survivors

```
median_age_survivors = data[data['Survived'] == 1]['Age'].median()
print(median_age_survivors)
```

8. Get the median age of non-survivors

```
median_age_non_survivors = data[data['Survived'] == 0]['Age'].median()
print(median_age_non_survivors)
```

9. Get the median fare of survivors

```
median_fare_survivors = data[data['Survived'] == 1]['Fare'].median()
print(median_fare_survivors)
```

10. Get the median fare of non-survivors

```
median_fare_non_survivors = data[data['Survived'] == 0]['Fare'].median()
print(median_fare_non_survivors)
```

Output:

The screenshot displays the output of a Jupyter Notebook. At the top, there are performance metrics: 'Average time' and 'Maximum time', both showing '0.113 s' and '113.00 ms'. A green status bar indicates '1 out of 1 shown test case(s) passed'. Below this, a 'Test case 1' section is shown with a green checkmark and a '113 ms' timer. The 'Expected output' and 'Actual output' columns are compared. The expected output shows a list of values: 233, 109, 468, 81, 217, 125, 427, 122, and 10.5. The actual output shows the same values, confirming the test case passed. The output is displayed in a table format with columns for 'Expected output' and 'Actual output'.

Expected output	Actual output
female --- 233	female --- 233
male --- 109	male --- 109
Name: Sex, dtype: int64	Name: Sex, dtype: int64
male --- 468	male --- 468
female --- 81	female --- 81
Name: Sex, dtype: int64	Name: Sex, dtype: int64
1 --- 217	1 --- 217
0 --- 125	0 --- 125
Name: Embarked_S, dtype: int64	Name: Embarked_S, dtype: int64
1 --- 427	1 --- 427
0 --- 122	0 --- 122
Name: Embarked_S, dtype: int64	Name: Embarked_S, dtype: int64
0.5398230088495575	0.5398230088495575
0.3810316139767055	0.3810316139767055
28.0	28.0
28.0	28.0
26.0	26.0
10.5	10.5

